

THREAT-MODEL

agentcore / ARI — Threat Model

1. What agentcore is (and isn't), security-wise
2. Attack surface
3. Findings (STRIDE)
4. Residual risk (accepted / host-owned)
5. ARI-Embedded security profile
6. Reporting

agentcore / ARI — Threat Model

Scope: the agentcore reference runtime and the security properties an ARI-conformant runtime should provide. **Method:** STRIDE over the real attack surface, grounded in the current code (src/core/, src/bindings/, python/agentcore/). **Status:** living document for a pre-1.0 project. This is a *static* review, not a substitute for an independent audit, fuzzing campaign, or penetration test of a deployed system.

1. What agentcore is (and isn't), security-wise

agentcore is a **library**, not a service or a sandbox. It runs an agent loop in your process. Two consequences drive everything below:

1. **Registered tool bodies execute with full host privileges.**
agentcore does not and cannot sandbox a tool you register. If you register a tool that calls `subprocess.run`, and a prompt-injected model invokes it, that is your design, not a runtime vulnerability.
2. **The dominant real-world risk is prompt-injection → tool-call.**
The model is untrusted output. The runtime's job is to *bound the blast radius* (caps, redaction, cancellation, backpressure) and to give you the primitives to apply least privilege — not to make arbitrary tools safe.

Trust boundaries

Zone	Examples	Treatment
Untrusted	LLM completions, tool-call args, tool results, message content/metadata from outside, a StateStore DB shared with other writers	Validate, cap, redact, never eval
Trusted (by configuration)	the host process, registered tool bodies, provider <code>base_url</code> , API keys, embedding application code	Out of scope to defend <i>from</i> ; in scope to avoid <i>leaking</i>

2. Attack surface

Surface	Entry point	Notes
---------	-------------	-------

Message ingress	Message.make, .content setter, router send	size caps, role set
Tool dispatch	ToolRegistry.invoke, ToolBox adapter	JSON in/out, size caps, error redaction
Provider egress	OpenAI/Anthropic/Ollama providers	network, API keys, base_url, timeouts
Multi-agent routing	AgentRouter inboxes	backpressure / memory
Persistence	InMemory/SQLiteStateStore	at-rest content, round-trip integrity
Concurrency	C++ shared_mutex, GIL release	data races, UB
Supply chain	build + optional deps	provenance, advisories

3. Findings (STRIDE)

Severity is qualitative (relative to a library at this stage). "Status" reflects this commit.

ID	STRIDE	Finding	Severity	
T1	DoS	Providers ignored timeout_ms/cancel_token; a hung provider call blocks an async worker thread indefinitely (8-thread pool → full stall).	Medium	Fixed — Agen OpenAI/Anth timeout and r
T2	Tampering / DoS	4 MiB Message.content cap enforced in make() but bypassable via direct .content = assignment.	Medium	Fixed — bind
T3	Info disclosure	Tool error text could leak secrets/paths; original redactor matched only a few prefixes.	Low-Med	Mitigated — (OpenAI/Anth URLs) + path returning exc strict for unt
T4	DoS	AgentRouter inboxes default to unbounded (max==0); a producer with no consumer exhausts memory.	Medium	Documented available; AR test asserts b
T5	Info disclosure	Conversation history is stored plaintext by SQLiteStateStore/in-memory.	Low	Documented responsibility StateStore); t
T6	DoS	RetryPolicy retried <i>all</i> exceptions, amplifying load on non-transient failures.	Low	Fixed — Valu guidance to n
T7	Elev. of privilege	No per-agent tool ACL: any agent in a Runtime can invoke any registered tool.	Low (design)	Documented future work.] domain.
T8	Spoofing / SSRF	Provider base_url (Ollama, OpenAI) is caller-trusted; an attacker-controlled URL exfiltrates prompts and API key.	Low-Med	Documented never from m

T9	Tampering	StateStore.load reads an arbitrary-size row fully into memory.	Low	Documented isolate it.
T10	Supply chain	No CodeQL / dependency-audit / SBOM / signed releases.	Medium (for a standard)	Partially addressed + pip-audit); recommended
T11	Info disclosure / integrity	Message.metadata was dropped on persistence load (lossy round-trip; trace metadata loss).	Low	Fixed — metadata

Memory safety & concurrency (C++ core)

- Access is guarded by `std::shared_mutex` consistently; `Engine::create_agent` documents its lock ordering to avoid a create/send race and deadlock.
- A ThreadSanitizer CI job exists (informational). **Recommended before 1.0:** an AddressSanitizer build in CI and a libFuzzer/atheris harness over tool-arg JSON parsing and message ingestion. Tracked as future work — not yet present.

4. Residual risk (accepted / host-owned)

- **Prompt-injection driving tool calls** — inherent to agents. Mitigate at the host: least-privilege tools, per-`Runtime` isolation, `ToolBox(strict=True)`, human-in-the-loop for dangerous actions. `agentcore` bounds blast radius; it does not eliminate this class.
- **Compromised dependency** in an optional provider extra — mitigated by audits
 - pinning; not eliminated.
- **At-rest disk compromise** — host's encryption responsibility.

5. ARI-Embedded security profile

For runtimes claiming **ARI-Embedded** (robots/edge/native), these are MUSTs, not SHOULDs, because the deployments are long-running, physically exposed, and often unattended:

1. **Bounded memory** — finite inbox bounds and fixed content/tool I/O caps by default (mitigates T4).
2. **Strict redaction** — strict-mode error handling by default (mitigates T3).
3. **Mandatory cancellation + wall-clock timeouts** on every provider/tool call (mitigates T1), thread-safe across the control loop.
4. **No hot-path allocation when observability is disabled** (the null trace sink).
5. **Declared, near-zero dependency footprint** (smaller attack surface).

6. Reporting

See [SECURITY.md](#) for private disclosure. Do not file public issues for vulnerabilities.